

Lending Club Loan Project

Danny Mathieson - March 2022

Downloading and Displaying the Dataset

```
In [2]: import numpy as np
import pandas as pd
```

```
In [3]: df = pd.read_csv('datasets/loan_data.csv')
df.head()
```

```
Out[3]:
```

	credit.policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333	10000.00
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000	10000.00
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000	10000.00
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333	10000.00
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000	10000.00

```
In [4]: df.shape
```

```
Out[4]: (9578, 14)
```

```
In [5]: # check for nulls
df.isna().sum()
```

```
Out[5]: credit.policy      0
purpose      0
int.rate     0
installment  0
log.annual.inc  0
dti          0
fico         0
days.with.cr.line  0
revol.bal    0
revol.util   0
inq.last.6mths  0
delinq.2yrs   0
pub.rec       0
not.fully.paid  0
dtype: int64
```

1. Feature Transformation - Transform Categorical Values into Numerical Values

```
In [6]: df.dtypes
```

```
Out [6]: credit.policy      int64
purpose      object
int.rate     float64
installment  float64
log.annual.inc float64
dti          float64
fico         int64
days.with.cr.line float64
revol.bal    int64
revol.util   float64
inq.last.6mths int64
delinq.2yrs  int64
pub.rec      int64
not.fully.paid int64
dtype: object
```

```
In [7]: # Only purpose needs to be changed to numerical values - get dummies
df_dummy = pd.get_dummies(df, drop_first=True)
df_dummy.head()
```

```
Out [7]:
```

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util
0	1	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1
1	1	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7
2	1	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6
3	1	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2
4	1	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5

```
In [8]: df_dummy.columns
```

```
Out [8]: Index(['credit.policy', 'int.rate', 'installment', 'log.annual.inc', 'dti',
               'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
               'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid',
               'purpose_credit_card', 'purpose_debt_consolidation',
               'purpose_educational', 'purpose_home_improvement',
               'purpose_major_purchase', 'purpose_small_business'],
              dtype='object')
```


2. EDA on Different Factors of the Dataset

```
In [9]: # Describe the data, split columns into either binary or numerical sub-types
df_dummy.describe()
numerical_cols = ['int.rate', 'installment', 'log.annual.inc', 'dti', 'fico', 'days.with.c
binary_cols = ['credit.policy', 'purpose_credit_card', 'purpose_debt_consolidation', 'p
```


```
In [10]: from matplotlib import pyplot as plt
import seaborn as sns
colors = plt.rcParams['axes.prop_cycle'].by_key()['color']
```

```
In [11]: # Loop over numerical columns - Box Plot overall and by fully paid status
for c in numerical_cols:
    # get datasets by category
    total_data = df_dummy[[c]]
    cat_data = df_dummy[[c, 'not.fully.paid']]
    # create boxplots
    sns.boxplot(x=None, y=c, data=total_data, ax=plt.subplot(1,2,1))
    sns.boxplot(x='not.fully.paid', y=c, data=cat_data, ax=plt.subplot(1,2,2))
    # format chart and show
    plt.suptitle(c)
    plt.xlabel('Not Paid Status')
    plt.show()
```


int.rate




installment




log.annual.inc




dti



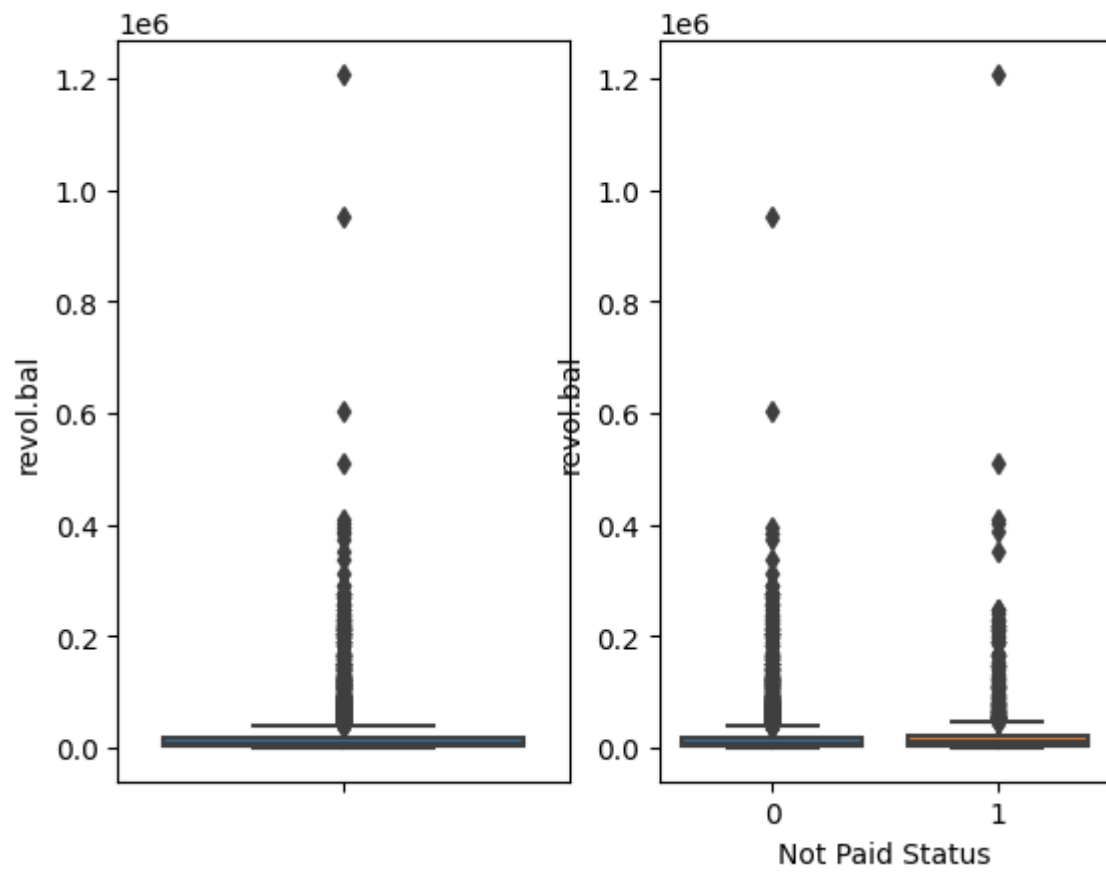
fico




days.with.cr.line




revol.bal




revol.util



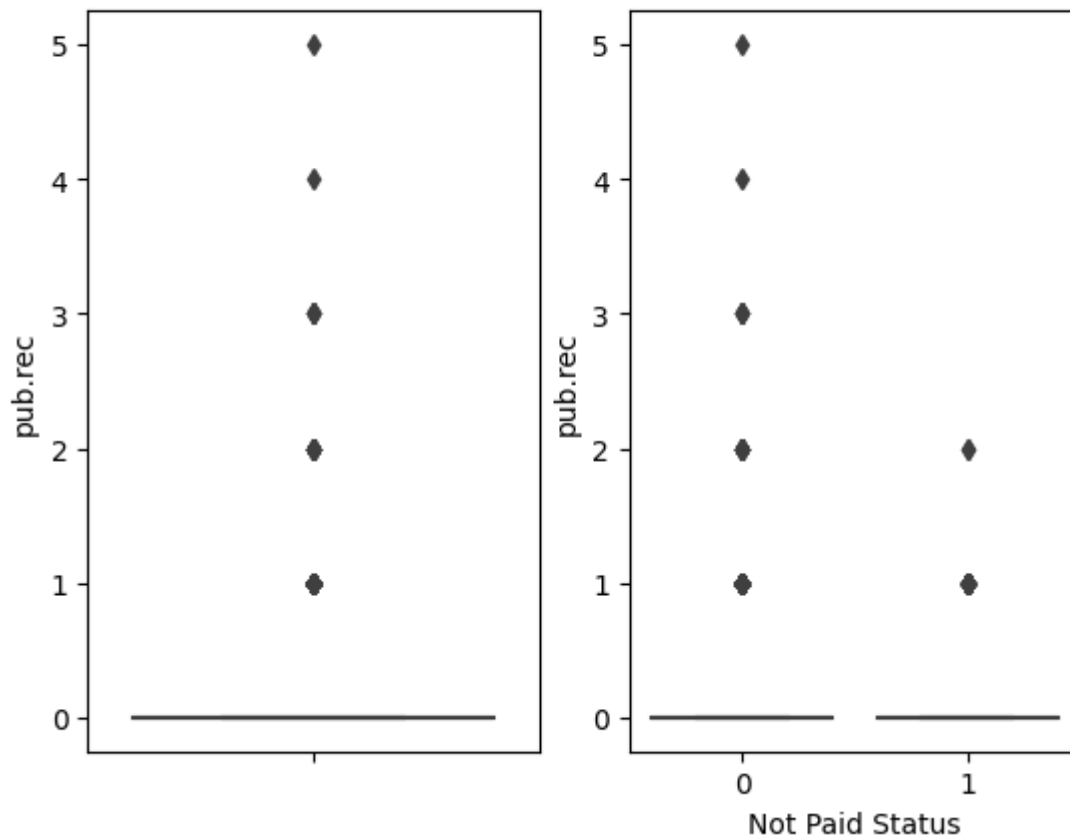
inq.last.6mths



delinq.2yrs



pub.rec



- int.rate

- Generally int.rate seems to correspond with a higher likelihood of not paying as the 25%, Median, and 75% are all higher than the paid class by a full percent or two
- The upper extreme doesn't seem to impact the outcome much as the max percent excluding outliers is the same and they seem to have even numbers of outliers on the high end
- interesting how no one with a sub 7% rate defaulted

- installment

- Once we start to get above 3.5 years the likelihood of default seems to go up, especially above 5 years

- log.annual.inc - nothing

- dti - nothing

- fico

- On first glance, lower fico scores definitely seem to have an impact on payback likelihood, but the entire range of paid back loans' fico scores contains the range of not paid back loans, including outliers.

- days.with.cr.line - nothing

- revol.bal - scale is too messed up to see much

- revol.util

- higher utilization rates has a slight impact on not paying back fully

- inq.last.6mths

- slightly impactful if above 2

- delinq.2yrs

- not enough non-zero data
- more outliers that have paid pack than haven't

pub.rec

- not enough non-zero data
- more outliers that have paid pack than haven't

```
In [12]: # Loop over Binary Columns & create a bar plot & stacked bar plot
for c in binary_cols:
    bin_data = df_dummy[[c, 'not.fully.paid']]
    bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
    bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
    # calculate percentages
    percent = bin_data.groupby(['Condition', 'Paid Status']).size().reset_index(name='count')
    percent['pct'] = percent['count'] / percent['count'].sum() * 100
    # create the plot
    order = {
        'Paid Status': ['No', 'Yes'],
        'Condition': ['No', 'Yes']
    }
    axis = sns.countplot(x='Paid Status', hue='Condition', data=bin_data, order=order)
    # add percentages for tooltips
    counter=0
    for p in axis.patches:
        h = p.get_height()
        pct = f"{round(percent['pct'][counter],1)}%"
        x_ax_pos = p.get_x() + p.get_width() / 2.0
        counter += 1
        axis.text(
            x_ax_pos,
            h + 3,
            pct,
            ha='center'
        )
    plt.title(c)
    plt.xlabel('Not Paid Status')
    plt.ylabel('Loans')
    plt.legend()
    plt.show()
```

/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

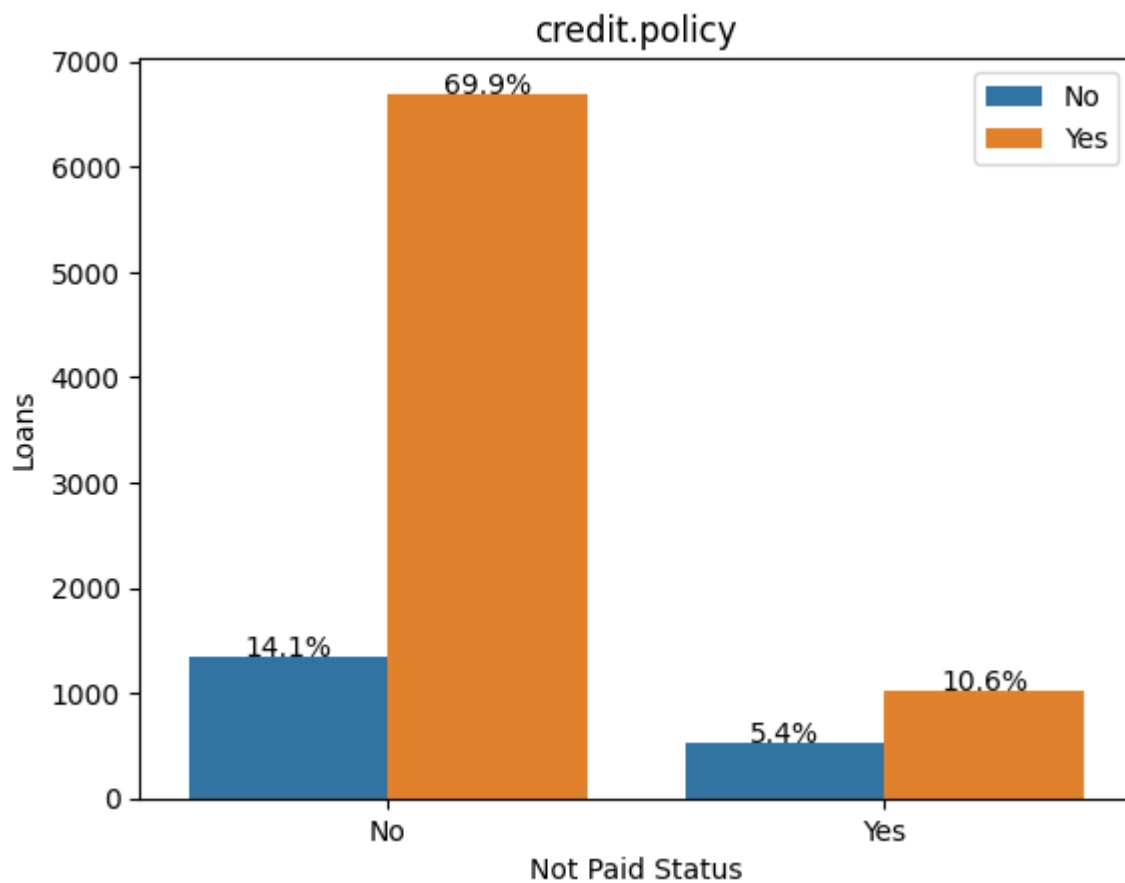
```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

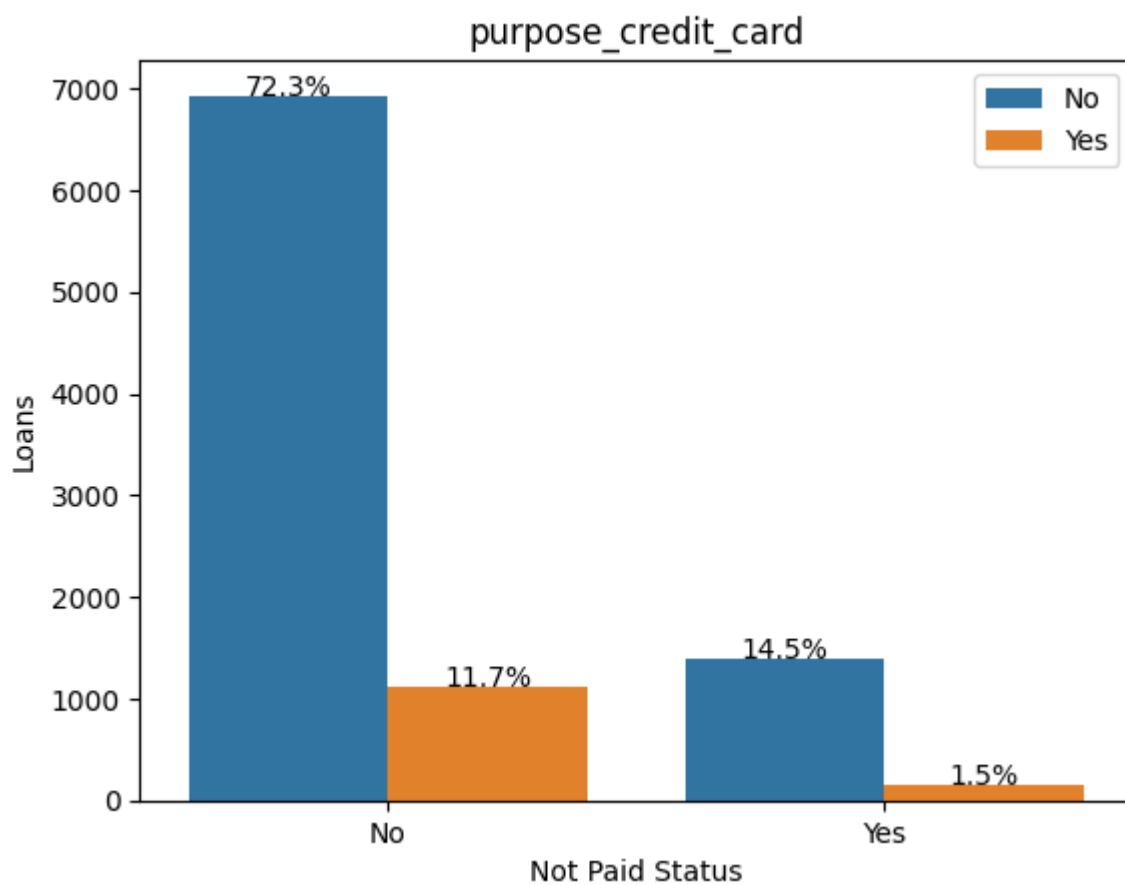
```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

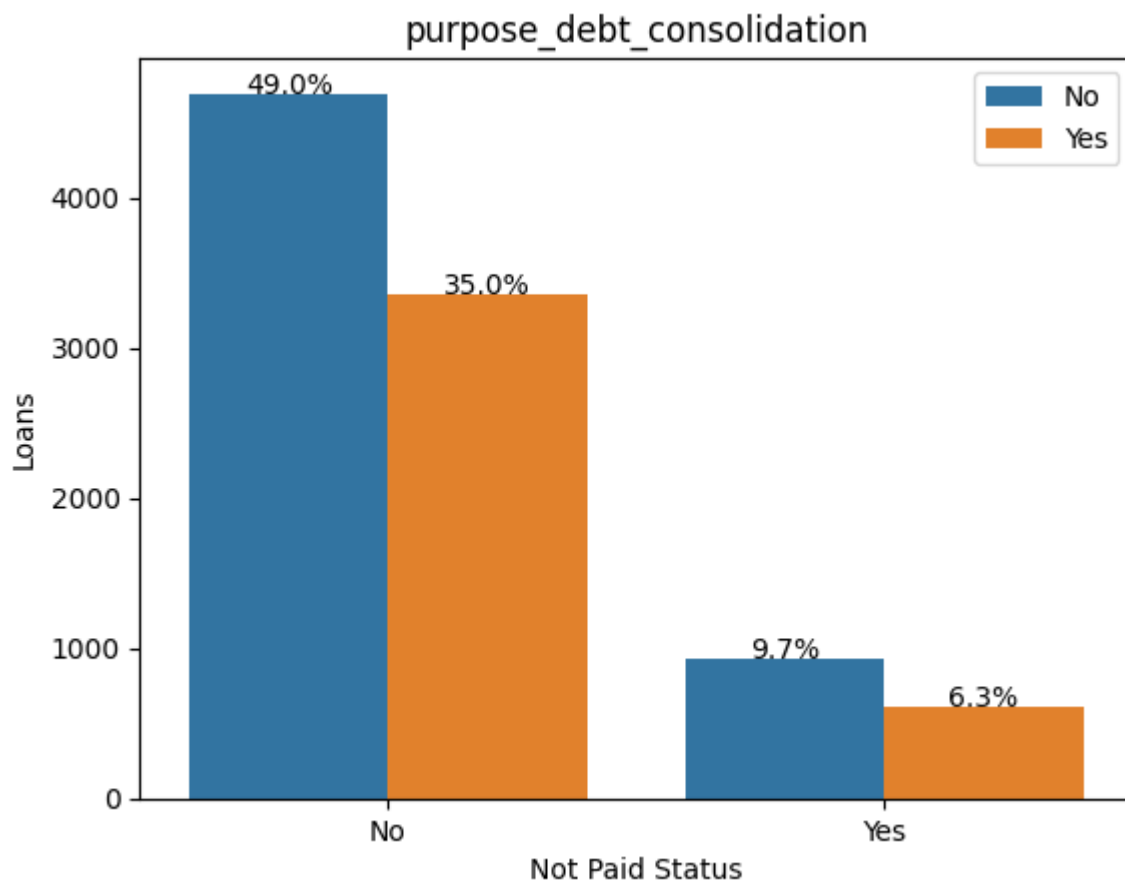
```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

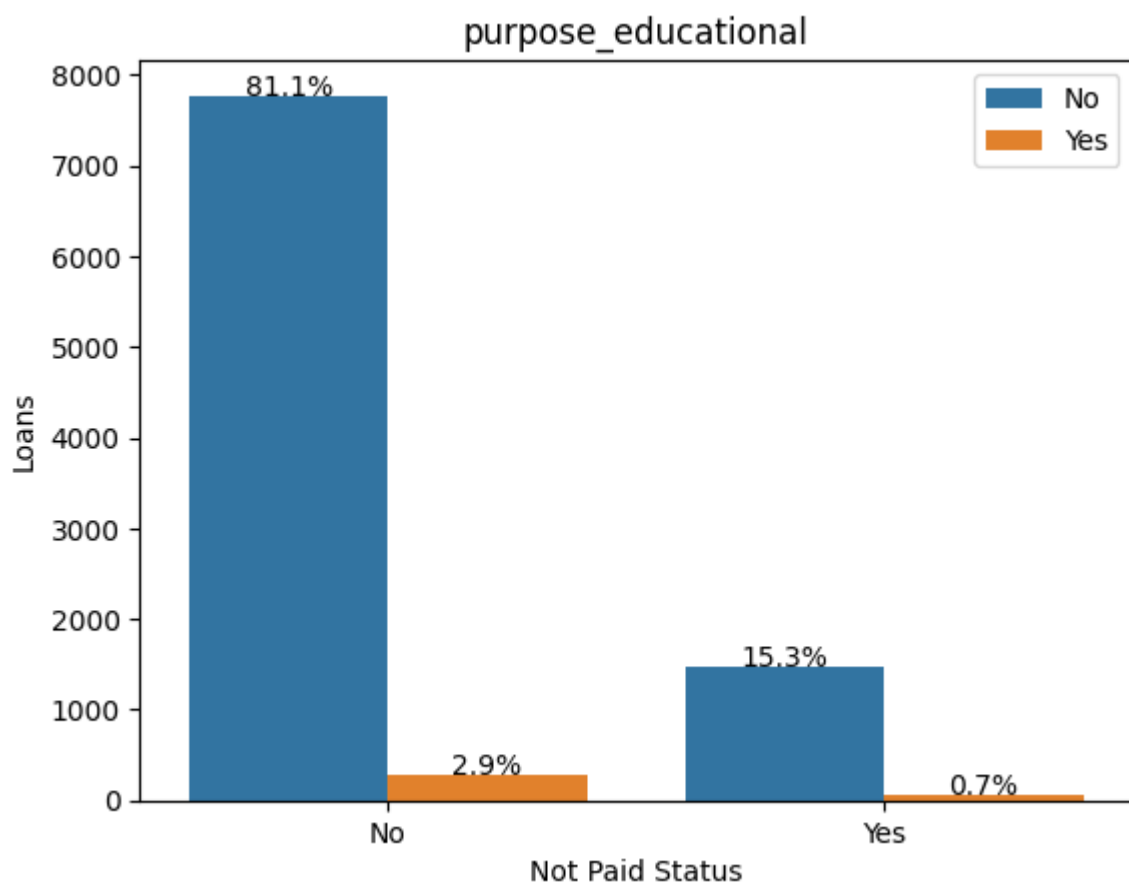
```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

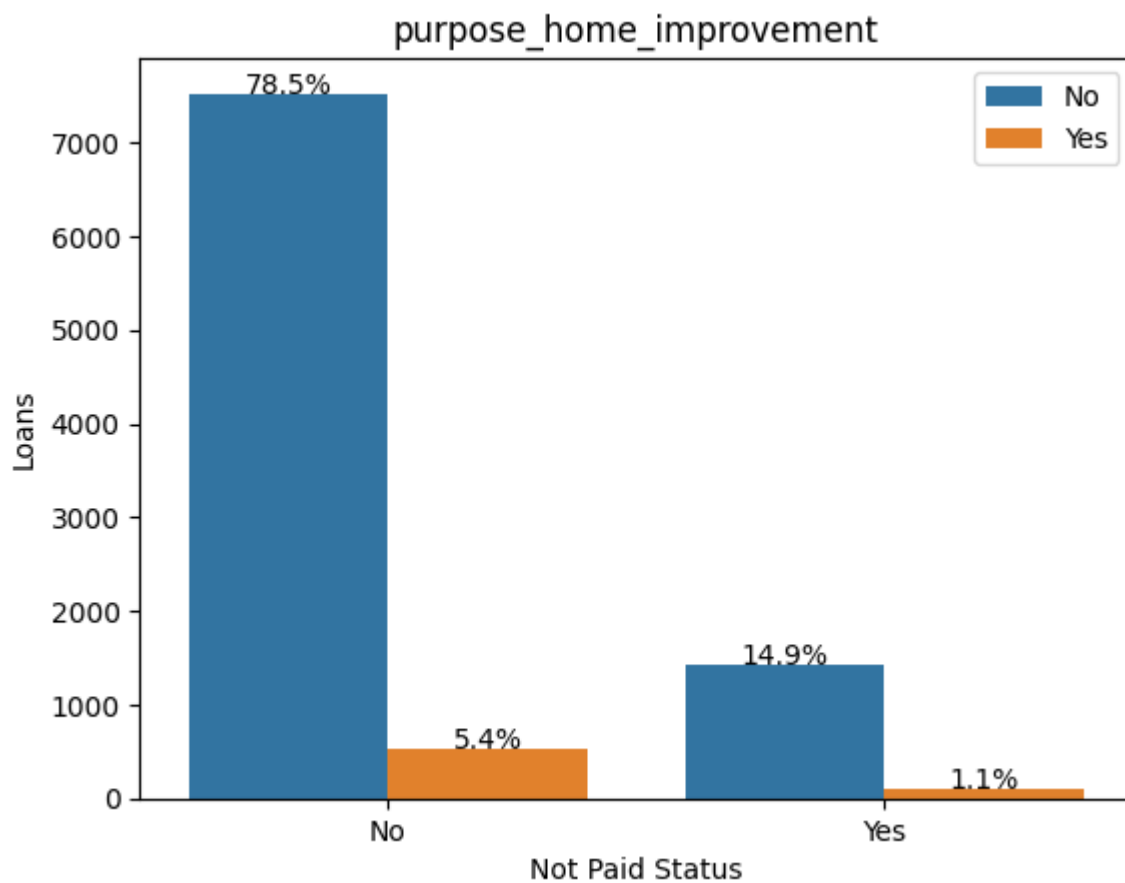
```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

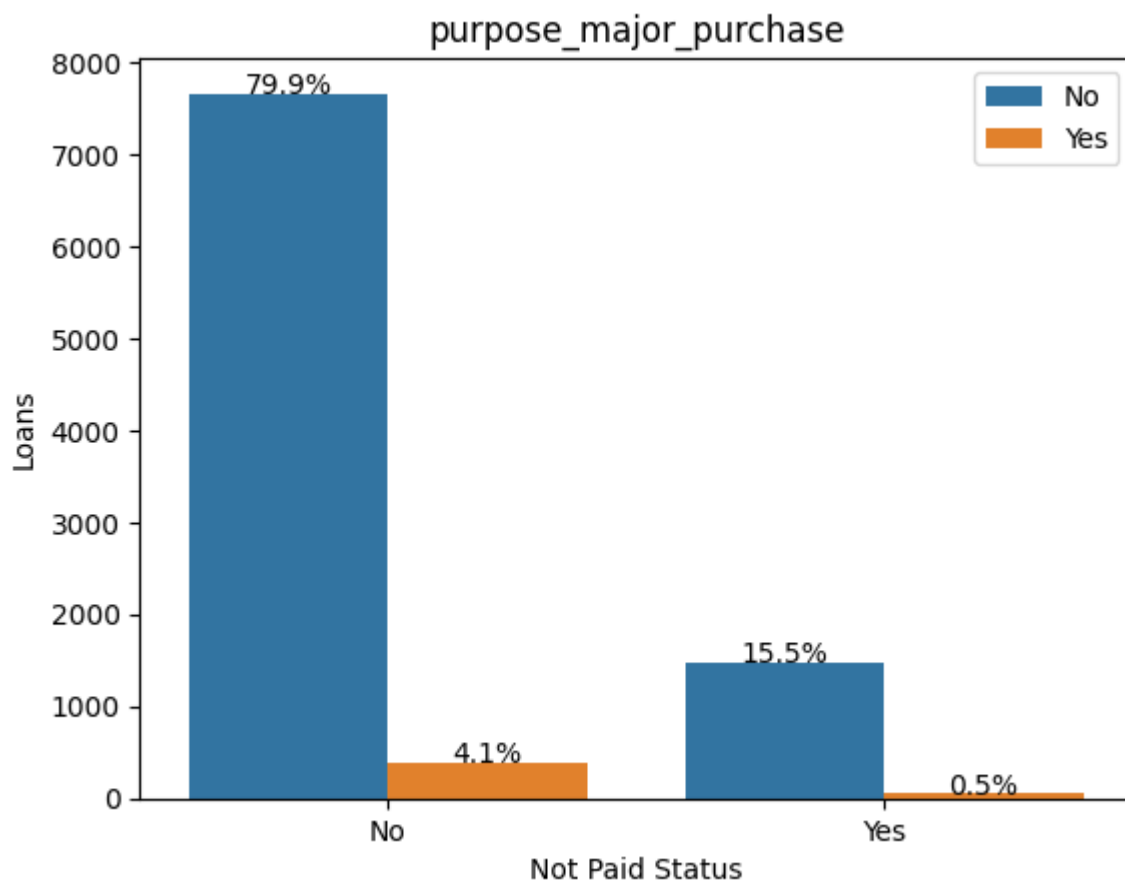
```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

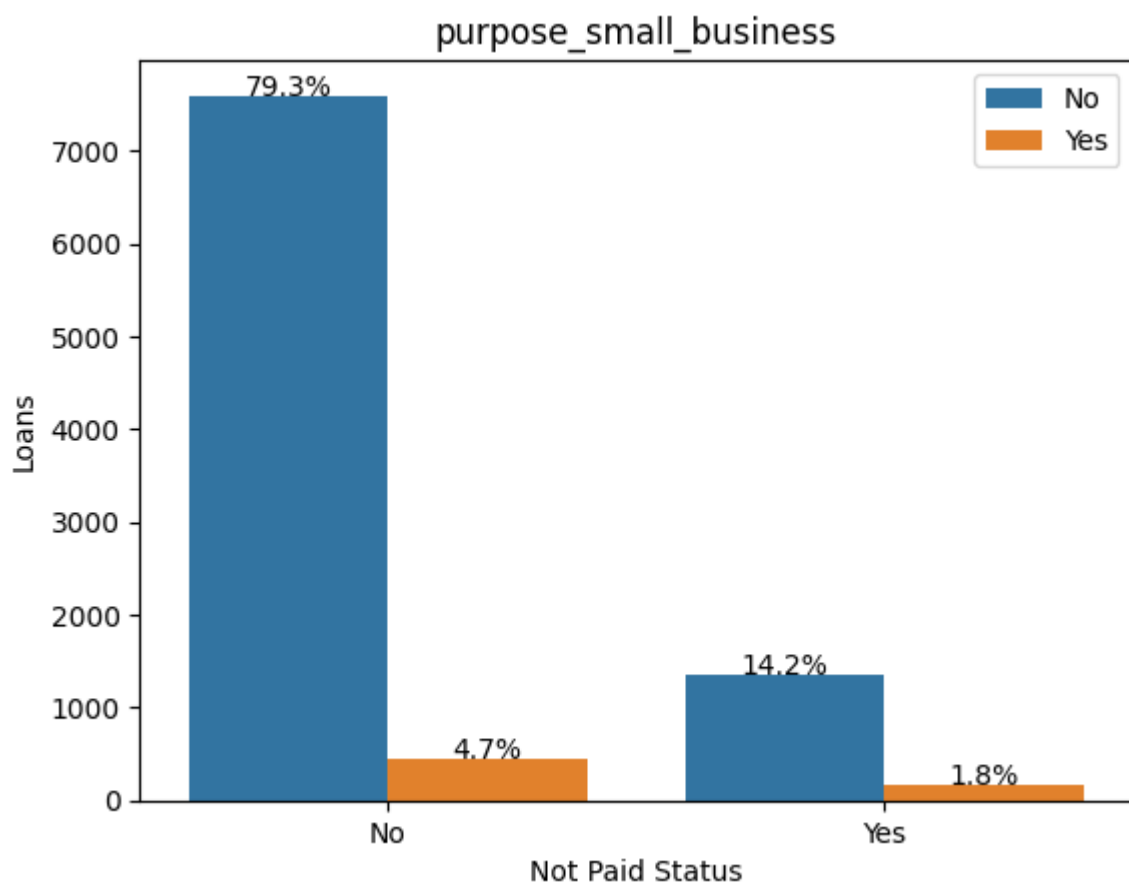
```
bin_data['Paid Status'] = np.where(df_dummy['not.fully.paid'] == 0, 'No', 'Yes')
```

```
/var/folders/r7/dtjny68152z02rjhb55rxv300000gn/T/ipykernel_94535/3147033967.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
bin_data['Condition'] = np.where(df_dummy[c] == 0, 'No', 'Yes')
```



- credit.policy

- About 4/5 of all people are approved for credit
- About 2/3 of the unpaid loans are from people that fit the credit policy criteria
- In general unapproved people are more likely to default. 25% of unapproved loans defaulted compared to just 13% of approved loans

- purpose_credit_card

- 13.2% of loans are credit card loans
- 11.4% of credit card loans default

- purpose_debt_consolidation

- 41.3% of loans are debt consolidation
- 15.2% of debt consolidation loans default

- purpose_educational

- Only 3.6% of loans are educational
- 19.4% of educational loans default

- purpose_home_improvement

- 6.5% of loans are home improvement
- 16.9% of home improvement loans default

- purpose_major_purchase

- 4.6% of loans are for a major purchase
- 10.9% of major purchase loans default

- purpose_small_business

- 6.5% of loans are for a small business
- 27.7% of small business loans default

- Key Takeaways


- Small Business loans are by far the riskiest, followed by educational
- Debt Consolidation is 41.3% of all loans, and they default at a slightly lower rate than average (16% default)
- Educational makes up such a small percentage, that small business is likely not super impactful
- The only other loan type that is above average is Home Improvement

3. Additional Feature Engineering

```
In [13]: # Find correlations between numerical features
features = df_dummy.drop(columns=['not.fully.paid'], axis=1)
num_features = features[numerical_cols]
labels = df_dummy[['not.fully.paid']]
```

```
In [14]: num_corr = num_features.corr()
sns.heatmap(data=num_corr, square=True, cmap='bwr')
```

Out[14]: <AxesSubplot: >



```
In [15]: corr_arr = num_corr.unstack()
corr_arr = corr_arr[corr_arr != 1]
corr_arr = corr_arr.drop_duplicates()
sorted_corr = corr_arr.sort_values(ascending=False)
opp_sorted_corr = corr_arr.sort_values(ascending=True)
```

```
print(f'Top Positive Correlations:\n\n{sorted_corr.head(10)}')
print(f'\n\nTop Negative Correlations:\n\n{opp_sorted_corr.head(10)}')
```

Top Positive Correlations:

int.rate	revol.util	0.464837
installment	log.annual.inc	0.448102
log.annual.inc	revol.bal	0.372140
dti	revol.util	0.337109
log.annual.inc	days.with.cr.line	0.336896
int.rate	installment	0.276140
fico	days.with.cr.line	0.263880
installment	revol.bal	0.233625
days.with.cr.line	revol.bal	0.229344
int.rate	dti	0.220006

dtype: float64

Top Negative Correlations:

int.rate	fico	-0.714821
fico	revol.util	-0.541289
dti	fico	-0.241191
fico	delinq.2yrs	-0.216340
	inq.last.6mths	-0.185293
	pub.rec	-0.147592
int.rate	days.with.cr.line	-0.124022
log.annual.inc	dti	-0.054065
revol.util	delinq.2yrs	-0.042740
days.with.cr.line	inq.last.6mths	-0.041736

dtype: float64

Correlation takeaways:

- The positive correlations are all below 0.5. I think we shouldn't remove any features due to those correlations
- FICO score is negatively correlated with quite a few features and should be removed because of it
- I may come back to check on int.rate & revol.util after building the model

```
In [16]: features = features.drop(columns=['fico'], axis=1)
```

4. Modeling

```
In [17]: # Data Preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, precision_recall_curve, confusion_matrix
import joblib

def data_split_standardise(x,y=None):
    if y is None:
        st=StandardScaler()
        st.fit(x)
        x_std=st.transform(x)
        joblib.dump(st,"model_objects/StandardScalar_trained.h5")
        return(x_std)
    else:
        x_train,x_test,y_train,y_test = train_test_split(x,y,random_state=0)
        st=StandardScaler()
        st.fit(x_train)
        x_train_std=st.transform(x_train)
        x_test_std=st.transform(x_test)
```

```
joblib.dump(st,"model_objects/StandardScalar_trained.h5")
return(x_train_std,x_test_std,y_train,y_test)
```

```
In [18]: x_train, x_test, y_train, y_test = data_split_standardise(features,labels)
```

```
In [19]: # Build Model
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Dense, Dropout
from tensorflow.keras.initializers import Constant
from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy, TruePositives
from tensorflow.keras.losses import BinaryCrossentropy
from tensorflow.keras.optimizers import Adam
from livelossplot import PlotLossesKerasTF
```

2023-03-12 10:43:27.803024: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [20]: # Modeling Constants
METRICS = [
    BinaryAccuracy(name='Binary_Accuracy'),
    Precision(name='Precision'),
    Recall(name='Recall'),
    TruePositives(name='True_Positives'),
    TrueNegatives(name='True_Negatives'),
    FalsePositives(name='False_Positives'),
    FalseNegatives(name='False_Negatives'),
    AUC(name='AUC'),
    AUC(name='Precision-Recall', curve='PR')
]
EPOCHS = 100
BATCH_SIZE = 512
```

2023-03-12 10:43:34.330481: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [21]: # Helpful plotting functions
```

```
# Confusion Matrix
def plot_cm(y_act, y_pred, p=0.5):
    cm = confusion_matrix(y_act, y_pred > p)
    plt.figure(figsize=(5,5))
    sns.heatmap(cm, annot=True)
    plt.title(f'Confusion Matrix P={p}')
    plt.ylabel('Actual')
    plt.xlabel('Predicted')

# ROC Curve
def plot_roc(y_act, y_pred, name='ROC', **kwargs):
    false_positive, true_positive, na = roc_curve(y_act, y_pred)

    plt.plot(false_positive, true_positive, label=name, **kwargs)
    plt.xlabel('False Positives')
    plt.ylabel('True Positives')
    plt.grid(True)
    ax = plt.gca()
    ax.set_aspect('equal')

# Precision-Recall Curve
def plot_prc(y_act, y_pred, name='ROC', **kwargs):
```

```
precision, recall, _ = precision_recall_curve(y_act, y_pred)

plt.plot(precision, recall, label=name, **kwargs)
plt.xlabel('Precision')
plt.ylabel('Recall')
plt.grid(True)
ax = plt.gca()
ax.set_aspect('equal')
```

```
In [22]: # First Attempt
def make_basic_model(metrics=METRICS, output_bias=None):
    if output_bias:
        output_bias = Constant(output_bias)



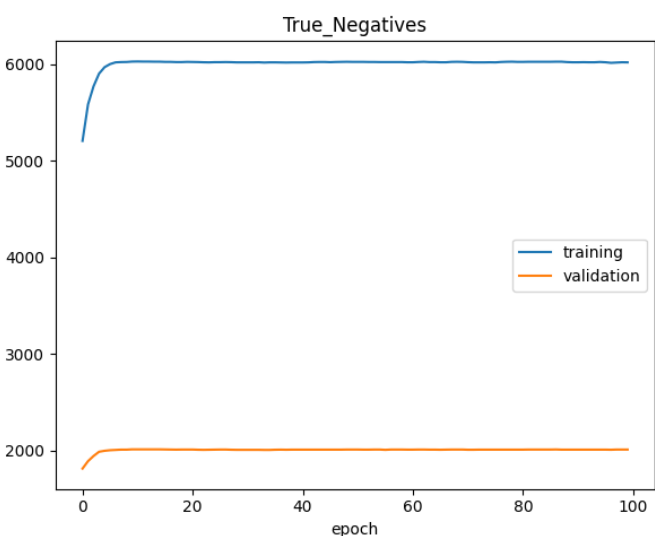





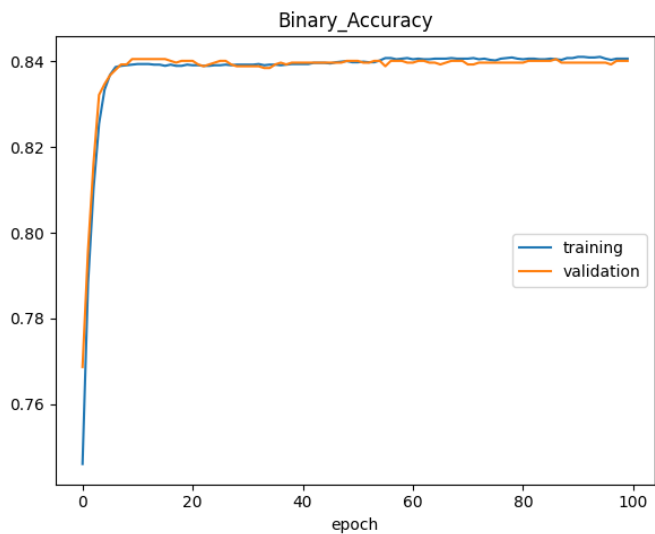

    model = Sequential()
    model.add(Input(shape=(None, x_train.shape[1]), name='Input_Layer'))
    model.add(Dense(12, activation='relu', name='Hidden_Layer_1'))
    model.add(Dense(8, activation='relu', name='Hidden_Layer_2'))
    model.add(Dense(1, activation='sigmoid', name='Output_Layer', bias_initializer=output_bias))


    model.compile(
        loss=BinaryCrossentropy(),
        optimizer=Adam(learning_rate=0.001),
        metrics=METRICS
    )

    return model
```

```
In [23]: model = make_basic_model(metrics=METRICS)
```

```
In [24]: model.fit(
    x_train,
    y_train,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    validation_data=(x_test, y_test),
    callbacks=[PlotLossesKerasTF()]
)
```





AUC

training	(min: 0.478, max: 0.702, cur: 0.702)
validation	(min: 0.459, max: 0.675, cur: 0.674)

Binary_Accuracy

training	(min: 0.746, max: 0.841, cur: 0.841)
validation	(min: 0.769, max: 0.841, cur: 0.840)

False_Negatives

training	(min: 999.000, max: 1151.000, cur: 1133.000)
validation	(min: 354.000, max: 381.000, cur: 380.000)

False_Positives

training	(min: 3.000, max: 825.000, cur: 12.000)
validation	(min: 1.000, max: 200.000, cur: 3.000)

Precision

training	(min: 0.149, max: 0.737, cur: 0.613)
validation	(min: 0.000, max: 0.333, cur: 0.250)

Precision-Recall

training	(min: 0.154, max: 0.314, cur: 0.314)
validation	(min: 0.141, max: 0.276, cur: 0.274)

Recall

training	(min: 0.001, max: 0.133, cur: 0.016)
validation	(min: 0.000, max: 0.071, cur: 0.003)

True_Negatives

training	(min: 5206.000, max: 6028.000, cur: 6019.000)
validation	(min: 1814.000, max: 2013.000, cur: 2011.000)

True_Positives

training	(min: 1.000, max: 153.000, cur: 19.000)
validation	(min: 0.000, max: 27.000, cur: 1.000)

Loss

training	(min: 0.404, max: 0.588, cur: 0.404)
validation	(min: 0.410, max: 0.562, cur: 0.411)

15/15 [=====] - 2s 126ms/step - loss: 0.4042 - Binary_Accuracy: 0.8406 - Precision: 0.6129 - Recall: 0.0165 - True_Positives: 19.0000 - True_Negatives: 6019.0000 - False_Positives: 12.0000 - False_Negatives: 1133.0000 - AUC: 0.7022 - Precision-Recall: 0.3136 - val_loss: 0.4110 - val_Binary_Accuracy: 0.8401 - val_Precision: 0.2500 - val_Recall: 0.0026 - val_True_Positives: 1.0000 - val_True_Negatives: 2011.0000 - val_False_Positives: 3.0000 - val_False_Negatives: 380.0000 - val_AUC: 0.6737 - val_Precision-Recall: 0.2744

Out[24]: <keras.callbacks.History at 0x13798f4f0>

```
In [25]: # Helper function for fitting & evaluating models
def evaluate_and_plot(model, x_train, x_test, y_train, y_test, batch_size=BATCH_SIZE):
    # Get Predictions & Evaluate
    train_preds = model.predict(x_train, batch_size=batch_size)
    test_preds = model.predict(x_test, batch_size=batch_size)
    results = model.evaluate(x_test, y_test, batch_size=batch_size, verbose=0)

    # Print Metric Scores
    print('\n\n')
    for metric, value in zip(model.metrics_names, results):
        print(f'{metric}:\t{value}')
```

```

# Plot Confusion Matrix
print('\n\n')
plot_cm(y_test, test_preds)
plt.show()

# Plot ROC
print('\n\n')
plot_roc(y_train, train_preds, name='Train', color=colors[0])
plot_roc(y_test, test_preds, name='Test', color=colors[1])
plt.legend()
plt.show()

# Plot Precision-Recall
print('\n\n')
plot_prc(y_train, train_preds, name='Train', color=colors[0])
plot_prc(y_test, test_preds, name='Test', color=colors[1])
plt.legend()
plt.show()

```

```

In [26]: evaluate_and_plot(
    model,
    x_train,
    x_test,
    y_train,
    y_test
)

```

WARNING:tensorflow:Model was constructed with shape (None, None, 17) for input KerasTensor(type_spec=TensorSpec(shape=(None, None, 17), dtype=tf.float32, name='Input_Layer'), name='Input_Layer', description="created by layer 'Input_Layer'"), but it was called on an input with incompatible shape (None, 17).


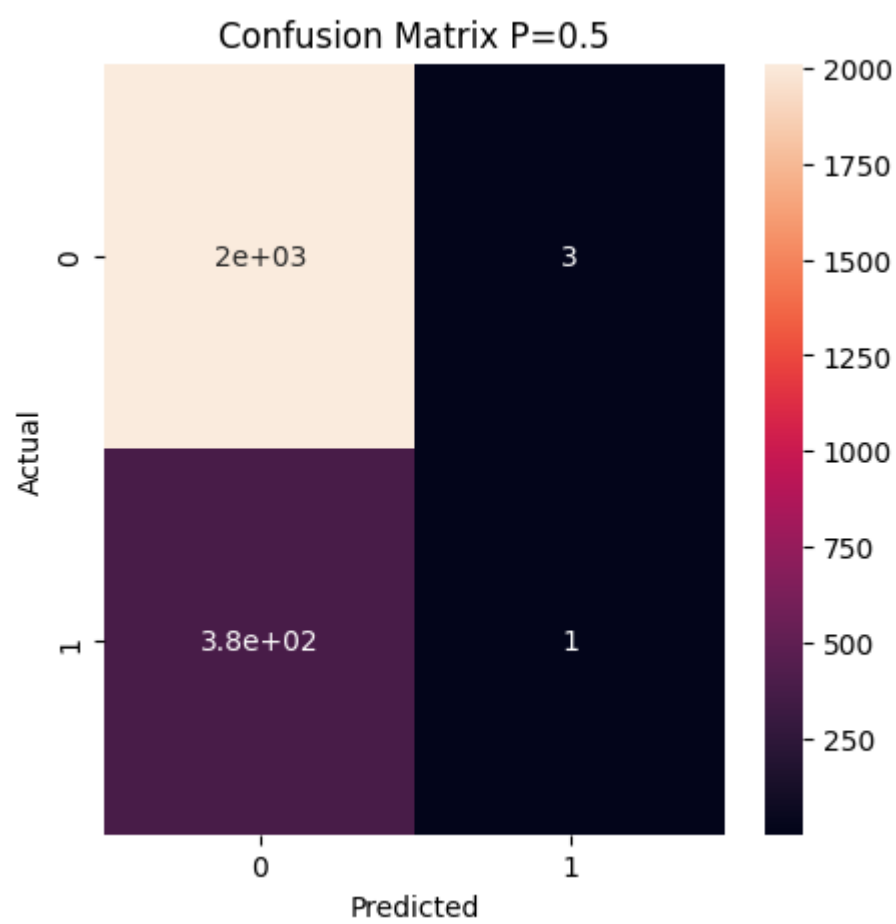
15/15 [=====] - 0s 885us/step


5/5 [=====] - 0s 1ms/step

```

loss:    0.4109605848789215
Binary_Accuracy:    0.8400834798812866
Precision:    0.25
Recall: 0.002624671906232834
True_Positives: 1.0
True_Negatives: 2011.0
False_Positives:    3.0
False_Negatives:    380.0
AUC:    0.6737378835678101
Precision-Recall:    0.27441197633743286

```





Model Overfit

Accuracy was great because we didn't predict any defaults - Class Imbalance











Let's try to add some initial bias



```
In [27]: # set initial bias
neg, pos = np.bincount(labels['not.fully.paid'])
initial_bias = np.log([pos/neg])
initial_bias
```

```
Out[27]: array([-1.65782418])
```

```
In [28]: model = make_basic_model(metrics=METRICS, output_bias=initial_bias)

model.fit(
    x_train,
    y_train,
    epochs=100,
    validation_data=(x_test,y_test),
    callbacks=[PlotLossesKerasTF()]
)
```





AUC

training	(min: 0.617, max: 0.746, cur: 0.745)
validation	(min: 0.638, max: 0.669, cur: 0.641)

Binary_Accuracy

training	(min: 0.839, max: 0.845, cur: 0.844)
validation	(min: 0.830, max: 0.839, cur: 0.833)

False_Negatives

training	(min: 1048.000, max: 1502.000, cur: 1053.000)
validation	(min: 355.000, max: 380.000, cur: 370.000)

False_Positives

training	(min: 26.000, max: 74.000, cur: 70.000)
validation	(min: 6.000, max: 38.000, cur: 29.000)

Precision

training	(min: 0.479, max: 0.646, cur: 0.586)
validation	(min: 0.111, max: 0.406, cur: 0.275)

Precision-Recall

training	(min: 0.252, max: 0.381, cur: 0.381)
validation	(min: 0.234, max: 0.267, cur: 0.238)

Recall

training	(min: 0.020, max: 0.090, cur: 0.086)
validation	(min: 0.003, max: 0.068, cur: 0.029)

True_Negatives

training	(min: 5957.000, max: 8018.000, cur: 5961.000)
validation	(min: 1976.000, max: 2008.000, cur: 1985.000)

True_Positives

training	(min: 29.000, max: 104.000, cur: 99.000)
validation	(min: 1.000, max: 26.000, cur: 11.000)

Loss

training	(min: 0.385, max: 0.431, cur: 0.385)
validation	(min: 0.414, max: 0.429, cur: 0.428)


225/225 [=====] - 2s 10ms/step - loss: 0.3853 - Binary_Accuracy: 0.8437 - Precision: 0.5858 - Recall: 0.0859 - True_Positives: 99.0000 - True_Negatives: 5961.0000 - False_Positives: 70.0000 - False_Negatives: 1053.0000 - AUC: 0.7445 - Precision-Recall: 0.3806 - val_loss: 0.4283 - val_Binary_Accuracy: 0.8334 - val_Precision: 0.2750 - val_Recall: 0.0289 - val_True_Positives: 11.0000 - val_True_Negatives: 1985.0000 - val_False_Positives: 29.0000 - val_False_Negatives: 370.0000 - val_AUC: 0.6409 - val_Precision-Recall: 0.2376



Out[28]: <keras.callbacks.History at 0x138593250>

```
In [29]: evaluate_and_plot(
    model,
    x_train,
    x_test,
    y_train,
    y_test
)
```

WARNING:tensorflow:Model was constructed with shape (None, None, 17) for input KerasTensor(type_spec=TensorSpec(shape=(None, None, 17), dtype=tf.float32, name='Input_Layer'), name='Input_Layer', description="created by layer 'Input_Layer'"), but it was called on an input with incompatible shape (None, 17).
15/15 [=====] - 0s 1ms/step
5/5 [=====] - 0s 1ms/step

loss: 0.4283325970172882
Binary_Accuracy: 0.8334029316902161
Precision: 0.2750000059604645
Recall: 0.028871390968561172
True_Positives: 11.0
True_Negatives: 1985.0
False_Positives: 29.0
False_Negatives: 370.0
AUC: 0.6409314274787903
Precision-Recall: 0.2376335859298706





Model Still Overfit - We actually Got Some Non-Payers Though!











Let's add some class weights to try to make sure we catch more fraud


```
In [30]: neg_weight = (1 / neg) * (labels.shape[0] / 2.0)
pos_weight = (1 / pos) * (labels.shape[0] / 2.0)
weights = {0: neg_weight, 1: pos_weight}
weights
```

```
Out[30]: {0: 0.5952765692977005, 1: 3.1239399869536855}
```

```
In [31]: model = make_basic_model(metrics=METRICS, output_bias=initial_bias)

model.fit(
    x_train,
    y_train,
    epochs=100,
    validation_data=(x_test,y_test),
    callbacks=[PlotLossesKerasTF()],
    class_weight=weights
)
```





AUC

training	(min: 0.569, max: 0.750, cur: 0.750)
validation	(min: 0.600, max: 0.672, cur: 0.651)

Binary_Accuracy

training	(min: 0.655, max: 0.799, cur: 0.662)
validation	(min: 0.599, max: 0.714, cur: 0.633)

False_Negatives

training	(min: 316.000, max: 1320.000, cur: 316.000)
validation	(min: 138.000, max: 249.000, cur: 165.000)

False_Positives

training	(min: 606.000, max: 2128.000, cur: 2112.000)
validation	(min: 435.000, max: 820.000, cur: 713.000)

Precision

training	(min: 0.242, max: 0.287, cur: 0.284)
validation	(min: 0.227, max: 0.253, cur: 0.233)

Precision-Recall

training	(min: 0.207, max: 0.355, cur: 0.353)
validation	(min: 0.230, max: 0.272, cur: 0.252)

Recall

training	(min: 0.139, max: 0.726, cur: 0.726)
validation	(min: 0.346, max: 0.638, cur: 0.567)

True_Negatives

training	(min: 3903.000, max: 7439.000, cur: 3919.000)
validation	(min: 1194.000, max: 1579.000, cur: 1301.000)

True_Positives

training	(min: 213.000, max: 836.000, cur: 836.000)
validation	(min: 132.000, max: 243.000, cur: 216.000)

Loss

training	(min: 0.589, max: 0.805, cur: 0.589)
validation	(min: 0.590, max: 0.684, cur: 0.639)


225/225 [=====] - 2s 10ms/step - loss: 0.5886 - Binary_Accuracy: 0.6620 - Precision: 0.2836 - Recall: 0.7257 - True_Positives: 836.0000 - True_Negatives: 3919.0000 - False_Positives: 2112.0000 - False_Negatives: 316.0000 - AUC: 0.7500 - Precision-Recall: 0.3534 - val_loss: 0.6390 - val_Binary_Accuracy: 0.6334 - val_Precision: 0.2325 - val_Recall: 0.5669 - val_True_Positives: 216.0000 - val_True_Negatives: 1301.0000 - val_False_Positives: 713.0000 - val_False_Negatives: 165.0000 - val_AUC: 0.6507 - val_Precision-Recall: 0.2521



Out[31]: <keras.callbacks.History at 0x139277310>

```
In [32]: evaluate_and_plot(
    model,
    x_train,
    x_test,
    y_train,
    y_test
)
```


WARNING:tensorflow:Model was constructed with shape (None, None, 17) for input KerasTensor(type_spec=TensorSpec(shape=(None, None, 17), dtype=tf.float32, name='Input_Layer'), name='Input_Layer', description="created by layer 'Input_Layer'"), but it was called on an input with incompatible shape (None, 17).
15/15 [=====] - 0s 1ms/step
5/5 [=====] - 0s 1ms/step

loss: 0.6390401124954224
Binary_Accuracy: 0.633402943611145
Precision: 0.2325080782175064
Recall: 0.5669291615486145
True_Positives: 216.0
True_Negatives: 1301.0
False_Positives: 713.0
False_Negatives: 165.0
AUC: 0.6507068872451782
Precision-Recall: 0.2520981729030609





Loss improved a bit, but still definitely overfit and in need of a bit more work


Additional Feature Engineering

Let's remove `log.annual.inc` and `revol.util` and look at correlations again

```
In [33]: non_feature_cols = ['not.fully.paid', 'fico', 'log.annual.inc', 'revol.util']
         features = df_dummy.drop(columns=non_feature_cols, axis=1)
         num_features = features.drop(columns=binary_cols, axis=1)
```

```
In [34]: num_corr = num_features.corr()  
sns.heatmap(data=num_corr, square=True, cmap='bwr')
```

Out[34]: <AxesSubplot: >



```
In [35]: corr_arr = num_corr.unstack()  
corr_arr = corr_arr[corr_arr != 1]  
corr_arr = corr_arr.drop_duplicates()  
sorted_corr = corr_arr.sort_values(ascending=False)  
opp_sorted_corr = corr_arr.sort_values(ascending=True)  
print(f'Top Positive Correlations:\n\n{sorted_corr.head(10)}')  
print(f'\n\nTop Negative Correlations:\n\n{opp_sorted_corr.head(10)}')
```

Top Positive Correlations:

int.rate	installment	0.276140
installment	revol.bal	0.233625
days.with.cr.line	revol.bal	0.229344
int.rate	dti	0.220006
	inq.last.6mths	0.202780
dti	revol.bal	0.188748
installment	days.with.cr.line	0.183297
int.rate	delinq.2yrs	0.156079
	pub.rec	0.098162
	revol.bal	0.092527

dtype: float64

Top Negative Correlations:











int.rate	days.with.cr.line	-0.124022
days.with.cr.line	inq.last.6mths	-0.041736
revol.bal	delinq.2yrs	-0.033243
installment	pub.rec	-0.032760
revol.bal	pub.rec	-0.031010
dti	delinq.2yrs	-0.021792
installment	inq.last.6mths	-0.010419
	delinq.2yrs	-0.004368
dti	pub.rec	0.006209
delinq.2yrs	pub.rec	0.009184



dtype: float64

```
In [36]: # Let's model with the same structure as the last run we did
x_train, x_test, y_train, y_test = data_split_standardise(features, labels)
```

```
In [37]: model = make_basic_model(metrics=METRICS, output_bias=initial_bias)

model.fit(
    x_train,
    y_train,
    epochs=100,
    validation_data=(x_test, y_test),
    callbacks=[PlotLossesKerasTF()],
    class_weight=weights
)
```





AUC

training	(min: 0.575, max: 0.718, cur: 0.718)
validation	(min: 0.572, max: 0.659, cur: 0.636)

Binary_Accuracy

training	(min: 0.635, max: 0.690, cur: 0.656)
validation	(min: 0.583, max: 0.676, cur: 0.639)

False_Negatives

training	(min: 379.000, max: 947.000, cur: 397.000)
validation	(min: 148.000, max: 244.000, cur: 176.000)

False_Positives

training	(min: 1581.000, max: 2220.000, cur: 2072.000)
validation	(min: 541.000, max: 850.000, cur: 688.000)

Precision

training	(min: 0.221, max: 0.269, cur: 0.267)
validation	(min: 0.202, max: 0.248, cur: 0.230)

Precision-Recall

training	(min: 0.213, max: 0.338, cur: 0.337)
validation	(min: 0.190, max: 0.264, cur: 0.241)

Recall

training	(min: 0.382, max: 0.671, cur: 0.655)
validation	(min: 0.360, max: 0.612, cur: 0.538)

True_Negatives

training	(min: 3811.000, max: 5983.000, cur: 3959.000)
validation	(min: 1164.000, max: 1473.000, cur: 1326.000)

True_Positives

training	(min: 509.000, max: 773.000, cur: 755.000)
validation	(min: 137.000, max: 233.000, cur: 205.000)

Loss

training	(min: 0.615, max: 0.742, cur: 0.615)
validation	(min: 0.620, max: 0.684, cur: 0.638)


225/225 [=====] - 2s 10ms/step - loss: 0.6151 - Binary_Accuracy: 0.6563 - Precision: 0.2671 - Recall: 0.6554 - True_Positives: 755.0000 - True_Negatives: 3959.0000 - False_Positives: 2072.0000 - False_Negatives: 397.0000 - AUC: 0.7175 - Precision-Recall: 0.3368 - val_loss: 0.6377 - val_Binary_Accuracy: 0.6392 - val_Precision: 0.2296 - val_Recall: 0.5381 - val_True_Positives: 205.0000 - val_True_Negatives: 1326.0000 - val_False_Positives: 688.0000 - val_False_Negatives: 176.0000 - val_AUC: 0.6358 - val_Precision-Recall: 0.2411



Out[37]: <keras.callbacks.History at 0x13b57feb0>

```
In [38]: evaluate_and_plot(
    model,
    x_train,
    x_test,
    y_train,
    y_test
)
```

WARNING:tensorflow:Model was constructed with shape (None, None, 15) for input KerasTensor(type_spec=TensorSpec(shape=(None, None, 15), dtype=tf.float32, name='Input_Layer'), name='Input_Layer', description="created by layer 'Input_Layer'"), but it was called on an input with incompatible shape (None, 15).
15/15 [=====] - 0s 872us/step
5/5 [=====] - 0s 1ms/step

loss: 0.6377151012420654
Binary_Accuracy: 0.6392484307289124
Precision: 0.22956326603889465
Recall: 0.5380577445030212
True_Positives: 205.0
True_Negatives: 1326.0
False_Positives: 688.0
False_Negatives: 176.0
AUC: 0.6357935667037964
Precision-Recall: 0.24106895923614502





Still not quite what we'd like... Seems like we're now predicting too many fraud cases.

Oversampling

Try oversampling on the positive class to make sure we're identifying potential loans that won't be repaid

```
In [58]: # Set up split datasets between pos and neg observations to sample at different rates
pos_df = df_dummy[df_dummy['not.fully.paid'] == 1].reset_index()
neg_df = df_dummy[df_dummy['not.fully.paid'] == 0].reset_index()
```



```
# pos_features = pos_df.drop(columns=non_feature_cols, axis=1)
# neg_features = neg_df.drop(columns=non_feature_cols, axis=1)

# pos_labels = pos_df['not.fully.paid']
# neg_labels = neg_df['not.fully.paid']
```

In [61]: pos_df.index

Out[61]: RangeIndex(start=0, stop=1533, step=1)

```
In [63]: # Randomly Sample the same number of positive observations as we have in the negative
resampled_pos_df = pos_df.sample(
    n=neg_df.shape[0],
    replace=True,
    random_state=0
)
resampled_pos_df.shape
```

Out[63]: (8045, 20)

```
In [68]: # re-combine the newly resampled dataset
resampled_df = pd.concat([resampled_pos_df, neg_df])

resampled_features = resampled_df.drop(columns=non_feature_cols, axis=1)
resampled_labels = resampled_df['not.fully.paid']

resampled_df.shape
```

Out[68]: (16090, 20)











```
In [69]: # confirm re-balanced classes
resampled_df['not.fully.paid'].value_counts()
```


Out[69]: 1 8045
0 8045
Name: not.fully.paid, dtype: int64

```
In [70]: # split the data into train & test
x_train, x_test, y_train, y_test = data_split_standardise(resampled_features, resample
```

```
In [72]: # re-build the model
# make sure to not add class weights or initial bias because we've rebalanced already
model = make_basic_model(metrics=METRICS, output_bias=initial_bias)

model.fit(
    x_train,
    y_train,
    epochs=100,
    validation_data=(x_test, y_test),
    callbacks=[PlotLossesKerasTF()]
)
```





AUC	training	(min: 0.546, max: 0.743, cur: 0.742)
	validation	(min: 0.679, max: 0.729, cur: 0.728)
Binary_Accuracy	training	(min: 0.565, max: 0.675, cur: 0.675)
	validation	(min: 0.629, max: 0.673, cur: 0.670)
False_Negatives	training	(min: 1923.000, max: 5799.000, cur: 1929.000)
	validation	(min: 461.000, max: 953.000, cur: 712.000)
False_Positives	training	(min: 1859.000, max: 2154.000, cur: 1992.000)
	validation	(min: 538.000, max: 922.000, cur: 614.000)
Precision	training	(min: 0.542, max: 0.675, cur: 0.675)
	validation	(min: 0.623, max: 0.675, cur: 0.675)
Precision-Recall	training	(min: 0.514, max: 0.731, cur: 0.731)
	validation	(min: 0.650, max: 0.699, cur: 0.697)
Recall	training	(min: 0.300, max: 0.683, cur: 0.682)
	validation	(min: 0.520, max: 0.768, cur: 0.641)
True_Negatives	training	(min: 3854.000, max: 7759.000, cur: 4016.000)
	validation	(min: 1115.000, max: 1499.000, cur: 1423.000)
True_Positives	training	(min: 2484.000, max: 4136.000, cur: 4130.000)
	validation	(min: 1033.000, max: 1525.000, cur: 1274.000)
Loss	training	(min: 0.593, max: 0.743, cur: 0.594)
	validation	(min: 0.608, max: 0.646, cur: 0.609)


378/378 [=====] - 3s 8ms/step - loss: 0.5936 - Binary_Accuracy: 0.6751 - Precision: 0.6746 - Recall: 0.6816 - True_Positives: 4130.0000 - True_Negatives: 4016.0000 - False_Positives: 1992.0000 - False_Negatives: 1929.0000 - AUC: 0.7424 - Precision-Recall: 0.7307 - val_loss: 0.6089 - val_Binary_Accuracy: 0.6704 - val_Precision: 0.6748 - val_Recall: 0.6415 - val_True_Positives: 1274.0000 - val_True_Negatives: 1423.0000 - val_False_Positives: 614.0000 - val_False_Negatives: 712.0000 - val_AUC: 0.7284 - val_Precision-Recall: 0.6974



Out [72]: <keras.callbacks.History at 0x140514220>

```
In [73]: evaluate_and_plot(
    model,
    x_train,
    x_test,
    y_train,
    y_test
)
```

WARNING:tensorflow:Model was constructed with shape (None, None, 16) for input KerasTensor(type_spec=TensorSpec(shape=(None, None, 16), dtype=tf.float32, name='Input_Layer'), name='Input_Layer', description="created by layer 'Input_Layer'"), but it was called on an input with incompatible shape (None, 16).
24/24 [=====] - 0s 1ms/step
8/8 [=====] - 0s 1ms/step

loss: 0.6089179515838623
Binary_Accuracy: 0.6703952550888062
Precision: 0.6747881174087524
Recall: 0.6414904594421387
True_Positives: 1274.0
True_Negatives: 1423.0
False_Positives: 614.0
False_Negatives: 712.0
AUC: 0.7283744215965271
Precision-Recall: 0.6973873376846313





Much better results!!